On Establishing a Benchmark for Evaluating Static Analysis Alert Prioritization and Classification Techniques

Sarah Heckman
North Carolina State University
Department of Computer Science
Campus Box 8206
Raleigh, NC, USA 27695-8206
sarah_heckman@ncsu.edu

Laurie Williams
North Carolina State University
Department of Computer Science
Campus Box 8206
Raleigh, NC, USA 27695-8206
williams@csc.ncsu.edu

ABSTRACT
Benchmarks provide an experimental basis for evaluating software engineering processes or techniques in an objective and repeatable manner. We present the FAULTBENCH v0.1 benchmark, as a contribution to current benchmark materials, for evaluation and comparison of techniques that prioritize and classify alerts generated by static analysis tools. Static analysis tools may generate an overwhelming number of alerts, the majority of which are likely to be false positives (FP). Two FP mitigation techniques, alert prioritization and classification, provide an ordering or classification of alerts, identifying those likely to be anomalies. We evaluate FAULTBENCH using three versions of a FP mitigation technique within the AWARE adaptive prioritization model. Individual FAULTBENCH subjects vary in their optimal FP mitigation techniques. Together, FAULTBENCH subjects provide a precise and general evaluation of FP mitigation techniques.

Categories and Subject Descriptors
D.2.4 [**Software Engineering**]: Software/Program Verification – *Reliability*, D.2.5 [**Software Engineering**]: Testing and Debugging – *Testing tools*

General Terms
Measurement, Reliability, Experimentation, Verification.

Keywords
Automated static analysis, alert prioritization, alert classification, benchmark creation, false positive mitigation

1. INTRODUCTION
Several open questions in software engineering involve evaluating processes and techniques that potentially improve aspects of the software development lifecycle. Empirical analysis of research theories are a component for acceptance of the theory within a research community [19]. Benchmarks provide an experimental basis for evaluating software engineering theories, represented by software engineering techniques, in an objective and repeatable manner [19]. A *benchmark* is defined as "a procedure, problem, or test that can be used to compare systems or components to each other or to a standard" [8]. Benchmarks represent the research problems of interest and solutions of importance in a research area through definition of the motivating comparison, task sample, and evaluation measures [18]. The task sample can contain programs, tests, and other artifacts dependent on the benchmark's motivating comparison. A benchmark controls the task sample reducing result variability, increasing repeatability, and providing a basis for comparison [18]. Additionally, successful benchmarks promote collaboration within a research community [18].

Several benchmarks in the realm of software anomaly[1] detection have emerged in recent years [15-17] containing subject programs of various sizes, in multiple languages, and with real or seeded faults. Current benchmarks provide meaningful points of comparison; however, they lack a detailed, repeatable process. Our goal is to supplement prior benchmarks by gathering a set of small, real, and anomalous Java programs from a variety of domains and providing a process for evaluation of the following software anomaly detection problem: how to identify which alerts generated by static analysis tools are program anomalies.

Static analysis tools can identify anomalies in source code early in the development process [8]. These tools produce reports listing possible program anomalies, which we call *alerts*. Inspection of each alert by a developer is required to determine if the alert is an indication of an important anomaly or a *true positive* (TP). When an alert is not an indication of an anomaly or is deemed unimportant to the developer (e.g. the alert indicates a programmer mistake inconsequential to program functionality), we call the alert a *false positive* (FP) [1]. Static analysis tools may generate an overwhelming number of alerts [11], the majority of which are likely to be FPs [6]. *Alert prioritization techniques*, used after static analysis is complete, can increase the usability of static analysis tools by presenting developers with TP alerts first. Additionally, *alert classification techniques*, used after static analysis is complete, can divide static analysis alerts into two groups: alerts likely to be TPs and alerts likely to be FPs. Prioritization and classification of static analysis alerts are both potential *FP mitigation techniques.*

The goal of our research is *to propose the FAULTBENCH benchmark to the software anomaly detection community for comparison and evaluation of FP mitigation techniques*. The literature in the realm of static analysis FP mitigation is moving towards a definition for conducting and evaluating research [10, 11, 13, 15, 21, 24]. FAULTBENCH provides a basis for comparison of static analysis FP mitigation techniques and contributes *subject programs; an analysis*

ESEM'08, October 9–10, 2008, Kaiserslautern, Germany.

[1] IEEE defines an anomaly "as a condition that deviates from expectations based on requirements specifications, design documents, user documents, or standards, or from someone's perceptions or experiences" [9]. The term anomaly encompasses software faults and other developer mistakes.

procedure; and *evaluation metrics*. The current version of FAULTBENCH, v0.1, contains six, open-source subject programs written in Java. We validate the selection of FAULTBENCH subject programs by comparing three versions of the AWARE [5] adaptive prioritization model (APM) FP mitigation techniques on static analysis alerts generated by the FINDBUGS [6] static analysis tool. FINDBUGS is a popular open source static analysis tool, which can identify faults like null pointers, open streams, etc. We describe how we created FAULTBENCH, and present the process of evaluating and comparing FP mitigation techniques. We invite researchers in the static analysis community to critique and improve the current benchmark.

The rest of the paper is structured as follows: Section 2 presents related work, Section 3 describes the FAULTBENCH creation and analysis procedure, Section 4 describes the FAULTBENCH case study, Section 5 presents the case study results, and Section 6 concludes and presents future work.

2. RELATED WORK

This section describes the related work in the areas of benchmark creation and the current static analysis FP mitigation techniques.

2.1 Current Anomaly Detection Benchmarks

There are several benchmarks in the realm of software anomaly detection. The SIEMENS [7] benchmark was created by researchers at Siemens Corporate Research and contains multiple versions of small C programs each containing a single anomaly and a suite of test cases. The benchmarks were created to evaluate control- and data-flow test adequacy criteria and were later used by Rothermel et al. [17] to evaluate regression test case prioritization.

BUGBENCH [15] is a benchmark containing seventeen buggy, open source, C/C++ applications ranging from seven thousand lines of code (KLOC) to 1028 KLOC in various domains. A Java benchmark for evaluation of the CHORD race condition detection static analysis tool [16] contains twelve concurrent programs ranging from 2.5 KLOC to 650 KLOC. PROMISE [3] is a repository for data sets from empirical research in predictive modeling, and half of the 60 data sets are for anomaly prediction. However, most of the PROMISE data sets provide metrics without the project source, and some data sets refer to large, open source projects while the remainders refer to commercial products. Other static analysis researchers [10, 11, 13, 21] have used large open source projects (e.g. Apache's httpd[2], Wine[3], Sun's JDK 1.6.0[4], Columba[5]) or commercial programs to evaluate FP mitigation techniques. While large open-source programs provide confidence and scale, the size of the sample evaluated (one to three programs) is a threat to external validity (e.g. the generalization of the results). Additional studies and subjects increase the generalization of experimental results [17]. Commercial examples show scalability of the technique in an industrial setting at the cost of repeatability and comparison.

These current benchmarks are insufficient for our needs for several reasons. First, current benchmarks are lacking a detailed, repeatable process for use and evaluation of FP mitigation techniques. Additionally, the current benchmarks are mostly for the C/C++ programming languages. Finally, alert prioritization research, especially adaptive prioritization, requires the removal of anomalies by a researcher unfamiliar with the program, which is costly for large projects with a large number of alerts. Therefore, we want to create a benchmark of relatively small, real, and anomalous Java programs from a variety of domains.

2.2 FP Mitigation Techniques

Kim and Ernst [10, 11] describe two static analysis alert prioritization techniques where the lifetime of a static analysis alert is measured from data mined from source code repositories. The lifetime of an alert is the time (in days) between alert creation and alert closure. One technique prioritizes static analysis alert types by the average lifetime of alerts sharing the same type [10]. Kim and Ernst assumed that alert types with shorter lifetimes have a higher ranking (e.g. alerts fixed quickly are likely important). However, alert types with shorter lifetimes could instead imply that those alert types are easiest to fix. The second technique weights alert types by the number of alerts closed by anomaly- and non-anomaly-fixes, where an anomaly-fix is a source code change where the developer fixes an anomaly or problem and a non-anomaly-fix is a change where an anomaly or problem is not fixed, like a feature addition [11]. The history based warning prioritization presented by Kim and Ernst [11] improves the alert precision by over 100% when compared to the alert precision of alerts prioritized by tool severity. However, the precision ranged from 17%-67%, which might be due to alert closures not having a causal relationship with the root cause of a anomaly-fix. Additionally, both prioritization techniques work best for fine-grained (e.g. many distinct alert types), homogeneous alert types. A homogeneous alert type means that all alerts sharing that type are either TPs or FPs. We utilize the idea of alert type homogeneity in our prioritization.

Williams and Hollingsworth [21] created a static analysis tool which evaluates how often the return values of method calls are checked in source code. A method is flagged with an alert when the return value for the method is inconsistently checked in calling methods. Williams and Hollingsworth use the HISTORYAWARE prioritization technique to prioritize methods by the percentage of time the return values of the methods are checked in the software repository and the current version of the code. The results show a statistically significant reduction of the FP rate when using the HISTORYAWARE prioritization technique on two case studies involving httpd[2] and Wine[3] applications.

Kremenek et al. [13] show that static analysis alerts in a similar alert locations tend to be homogeneous. On average, 88% of methods, 52% of files, and 13% of the directories with two or more alerts contained homogeneous alerts. Kremenek et al. created a FEEDBACK-RANK algorithm whereby the developer's feedback is used to prioritize the remaining alerts. The static analysis tools used by Kremenek et al. take advantage of understanding where a static analysis tool checked for an alert, but did not find a potential anomaly [14]. We also use the developer's feedback to drive the adaptive prioritization of un-inspected static analysis alerts, and Kremenek et al inspired our version of the code locality alert characteristic.

Boogerd and Moonen [4] present the ELAN technique to prioritize static analysis alerts by their execution likelihood, which is "the probability that a given program point will be executed at least once in an arbitrary program run." The prioritization is a measure of alert severity relative to the program under analysis. While the results showed that the prioritization technique did prioritize alerts by execution likelihood by comparison with unit test coverage, the

[2] http://httpd.apache.org/
[3] http://www.winehq.org/
[4] http://java.sun.com/javase/
[5] http://columba.sourceforge.net/

analysis did not investigate if the prioritization identified more alerts of interest to the developer.

3. BENCHMARK CREATION

The goal of contributing the FAULTBENCH benchmark is to create a (1) *suite of subject programs and alert oracles* and (2) *repeatable procedures for evaluation of FP mitigation techniques.* We have created a benchmark of Java programs from various domains, ranging from 1,276 – 14,120 lines of code (LOC) and static analysis alert oracles from alerts generated by FINDBUGS [6]. FINDBUGS uses code scans, control-flow, and data-flow analysis to detect common source code patterns that are possible anomalies [6]. FINDBUGS [6] detects 331 distinct alert types at three priority levels. We used the FINDBUGS Eclipse[6] plug-in [6] to generate alerts on the subject programs with-in the Eclipse workbench. We configured FINDBUGS to report alerts at all priority and effort levels, which maximizes alerts reported by FINDBUGS.

Below, we define the process for evaluating adaptive FP mitigation techniques to provide motivation for the creation of FAULTBENCH. We then define the purpose and describe how FAULTBENCH fulfills properties for successful benchmarks. In addition, we provide the steps for choosing benchmark subjects and initializing those subjects for use in FAULTBENCH.

3.1 FAULTBENCH Process

We present the steps for evaluating adaptive FP mitigation techniques with FAULTBENCH. Non-adaptive FP mitigation techniques would only need to evaluate the prioritized or classified alerts without fixing or suppressing alerts. For adaptive FP mitigation techniques, the states of the alerts are recorded after each inspection. An alert may be in one of three states: uninspected, TP, or FP. The FAULTBENCH process is as follows:

1. Run a static analysis tool against a clean version of the subject program. (If the static analysis can run automatically, turn on the option.)
2. Record the original state of the alert set.
3. Prioritize or classify the generated alerts with a FP mitigation technique.
4. Starting at the top of the prioritized list or randomly selecting an alert classified as important, inspect each alert,
   a. If the alert oracle indicates the alert is an anomaly, then fix the alert with the specified change. If the static analysis tool does not run automatically, then rerun static analysis.
   b. If the alert oracle indicates the alert is a FP, then suppress the alert.
5. After each alert inspection, record the state of the alert set.
6. After all alert inspections, evaluate the results via the evaluation metrics provided in Section 3.2.3.

3.2 Definition of FAULTBENCH

We define FAULTBENCH in terms of the three components presented by Sim et al. [18]: motivating comparison, task sample, and evaluation measures.

3.2.1 Motivating Comparison

The motivating comparison advocated by Sim et al. [18] describes why the results of comparing two tools or techniques are important for furthering the research surrounding the comparison. The motivating comparison of FAULTBENCH is to find the static analysis FP mitigation technique with the best rate of anomaly detection. Static analysis is an effective means of anomaly removal [23] and is cost effective with the detection of three to four potential field failures [20]. However, a large number of alerts, especially FP alerts, leads to rejection of the tool [4]. Specifically, we can use FAULTBENCH to answer the following research questions:

- [Q1]: Can alert prioritization improve the rate of anomaly detection when compared to the tool's output?
- [Q2]: How does the rate of anomaly detection compare between alert prioritization techniques?
- [Q3]: Can alert categorization correctly predict TP and FP alerts?

3.2.2 Task Sample

The task sample is a representative sample of tests that FP mitigation techniques should solve [18]. For FAULTBENCH, the task sample consists of (1) six real Java subject programs ranging from 1,276 – 14,120 lines of code (LOC); (2) the set of FINDBUGS [6] alerts identified as TP or FP in the context of the subject programs (*alert oracle)*; (3) a set of source code changes to fix each TP alert; and (4) the experimental control alert prioritizations: OPTIMAL, TOOL, and RANDOM. Section 3.4 describes the subject program selection process for FAULTBENCH. The descriptions for creating the remaining task sample data is in Section 3.5, FAULTBENCH Initialization.

3.2.3 Evaluation Measures

FAULTBENCH evaluates static analysis FP mitigation techniques. *Alert prioritization techniques* order alerts such that alerts likely to be indications of important anomalies are at the top of an alert list. *Alert classification techniques* divide static analysis alerts into two groups: alerts likely to be TPs and alerts likely to be FPs. Alert prioritization classifies alerts when alerts are ranked on a divisible numerical scale.

Alert prioritization evaluation uses the *Spearman rank correlation*, which evaluates alert orderings by measuring the distance between the rank, or location, of the same alert between two orderings. Users of the benchmark compare alert prioritization generated by a prioritization technique with an OPTIMAL ordering of alerts. An alert prioritization highly correlated with OPTIMAL at a statistically significant level suggests that the prioritization technique correctly ordered alerts such that alerts likely to indicate anomalies are higher in an alert list.

Alert classification techniques predict if alerts are TPs or FPs. If we classify an alert as a TP when the alert is a TP, then we have correctly classified the alert and we call that classification a *true positive classification (TP_C).* Additionally, if we classify an alert as a FP when the alert is not an indication of an anomaly we have correctly classified a negative prediction, which we call a *true negative classification (TN_C)*. A *false positive classification (FP_C)* is when the model predicts that an alert is a TP (a positive classification) when the alert is actually not an indication of an anomaly. A *false negative classification (FN_C)* is when the model predicts that an alert is a FP (a negative classification) when the alert is actually an anomaly. We are focusing on the classification of alerts identified by the static analysis tool; therefore, we are not considering software anomalies not found by static analysis tools. Figure 1 is a classification table.

[6] Eclipse is an open source integrated development environment. Eclipse may be found at: http://eclipse.org

Figure 1: Classification Table (adapted from Zimmerman et al. [24] where they used fault where we use anomaly)

		Anomalies are observed. True	Anomalies are observed. False	
Model predicts alerts	Positive	True Positive (TP_C)	False Positive (FP_C)	**Precision**
	Negative	False Negative (FN_C)	True Negative (TN_C)	
		Recall		**Accuracy**

The following metrics [11, 21, 22, 24] are used to evaluate the classification of static analysis alerts:

- **Precision**: the proportion of correctly classified anomalies (TP_C) out of all alerts predicted as anomalies (TP_C + FP_C). The precision calculation is presented in Equation 1.

$$precision = \frac{TP_C}{TP_C + FP_C} \quad (1)$$

- **Recall**: the proportion of correctly classified anomalies (TP_C) out of all possible anomalies (TP_C + FN_C). The recall calculation is presented in Equation 2.

$$recall = \frac{TP_C}{TP_C + FN_C} \quad (2)$$

- **Accuracy**: the proportion of correct classifications out of all classifications. The accuracy calculation is presented in Equation 3

$$accuracy = \frac{TP_C + TN_C}{TP_C + TN_C + FP_C + FN_C} \quad (3)$$

- **Anomaly Detection Rate Curve**: the area under the curve of the cumulative percentage of anomalies detected after each inspection. An example is in Figure 5.

3.3 Desiderata for Benchmarks

Sim et al. [18] describe seven properties of successful benchmarks: accessibility, affordability, clarity, relevance, solvability, portability, and scalability. Lu et al. [15] also provide five benchmark selection criteria: representative, diverse, portability, accessibility, and fairness. The following subsections describe how FAULTBENCH meets these desiderata:

- **Accessibility:** A benchmark should be easy to obtain and use. Each of the FAULTBENCH subjects is available online through various open source licenses. The subject programs, generated alerts, anomaly fixes, and evaluation materials related to FAULTBENCH are publicly available at http://agile.csc.ncsu.edu/faultbench.
- **Affordability:** A benchmark's cost (e.g. human, software, and hardware resources) should be comparable to the value of the results. To complete the benchmark for a single prioritization technique takes 8-10 hours on a single computer. Additional time is required for evaluating further techniques.
- **Clarity:** A benchmark's documentation should be clear and concise. The FAULTBENCH documentation is provided at http://agile.csc.ncsu.edu/faultbench for evaluation and comparison of other FP mitigation techniques to ensure repeatability and disclosure.
- **Relevance/Representative:** A benchmark must contain representative subjects and performance measures related to the motivating comparison. FAULTBENCH contains Java programs from various domains created by developers of varying levels of experience. The performance measures are standard in the area of data mining [22], software anomaly detection [24], and static analysis FP mitigation [21].
- **Solvability:** Completing the task sample and obtaining correct metrics is not difficult. The task samples vary in size and number of FINDBUGS static analysis alerts. FINDBUGS identified 55 alert types in the task sample from the 331 possible alert types. Additionally, an analysis program is provided as part of the benchmark materials
- **Portability:** A benchmark should be useable by different FP mitigation techniques without bias. The task sample consists of stand-alone Java projects containing required libraries. Use of the Java language assumes platform portability.
- **Scalability/Fairness:** A benchmark should be scalable to varying FP mitigation techniques and not have bias towards a specific technique. Currently, FAULTBENCH contains Java subject programs and can only evaluate FP mitigation techniques on alerts generated by Java static analysis tools. FAULTBENCH supplements other benchmarks in C and C++ and the authors encourage expansion to the benchmark as demonstrated in Section 3.2.

3.4 FAULTBENCH Subject Selection

The goal of FAULTBENCH is to provide a benchmark to the software anomaly detection community for comparison and evaluation of static analysis FP mitigation techniques. Therefore, the subject programs in the benchmark must meet the following criteria: open source; small (less than 15 KLOC), of various domains, written in Java; and compliable with Java 1.4.2 or Java 1.5. To find possible subject programs, we investigated the benchmarks presented in the related work section. None of the subjects in those benchmarks met our criteria for selection. Next, we investigated programs analyzed by the static analysis, style checker tool PMD[7]. The PMD website maintains a page reporting results from running PMD on SourceForge[8] projects. We investigated the 15 smallest programs (based on the number of analyzed non-commented source LOC) for inclusion into our benchmark and selected 11 as possible subjects. The un-chosen projects did not contain source releases or no longer existed as projects. Then, we searched for small components of commonly used libraries and applications, like Apache and Eclipse. One subject was identified when satisfying the library requirements of an earlier subject. The final potential subject comes from a student project associated with the authors' research group. Table 1 presents the set of possible subject programs.

The set of subjects were further refined through an analysis of six characteristics: domain; number of developers; LOC; number of FINDBUGS alerts; maturity; and alert distribution. First, we quantified each of the characteristics. For the categorical characteristics (e.g. domain and maturity), we assigned a numerical value to each category. The alert distribution is a value describing how many unique alert types FINDBUGS identified in a subject program. The alert distribution is the sum of the number of alerts of the same type in a subject divided by the number of subjects that contain at least one alert of the that type which is then divided by

[7] http://pmd.sourceforge.net/

[8] SourceForge is a repository for open source projects: http://sourceforge.net

Table 1: Potential FAULTBENCH benchmark subjects

Subject	Version	License	Domain	# Dev	# LoC	# Alerts	Maturity	Alert Dist.	Area
commons	2005.05.30	GNU LGPL	1 - software dev	2	5560	70	5 - Production	0.38	173,497.0
commons-logging	1.1.1	Apache 2.0	1 - software dev	12	5426	126	5 - Production	0.34	324,513.6
csvobjects	0.5beta	GNU GPL	2 - data format	1	1577	7	5 - Production	0.64	5,477.5
importscrubber	1.4.3	Apache Software License	1 - software dev	2	1653	35	4 - Beta	0.31	26,545.7
itrust	Fall 2007	Educational	3 - web	5	14120	110	3 - Alpha	0.61	703,277.0
javaserver	5.1	Artistic	6 -communication	1	1752	31	5 - Production	0.39	24,348.0
jbook	1.4	GNU GPL	7 - educational	1	1276	52	5 - Production	0.28	29,400.9
jdom	1.1	Apache-style	2 - data format	3	8422	55	5 - Production	0.19	211,638.6
junit-addons	1.4	Apache Software License	1 - software dev	1	4856	109	4 - Beta	0.45	231,488.3
kaprekar	3.0	GNU GPL, MPL 1.1	5 - math	1	1869	33	4 - Beta	0.21	27,576.4
mflow	0.1	GNU GPL	6 - communication	1	4172	86	3 - Alpha	0.33	157,283.6
org.eclipse.core.runtime	3.3.1.1	Eclipse Public License	1 - software dev	100	2791	98	5 - Production	0.30	239,546.9
schemalizer	0.16	GNU LGPL	2 - data format	1	2524	29	3 - Alpha	0.17	32,826.6
xmlwriter	2.2.2	BSD License	2 - data format	2	953	6	5 - Production	0.70	3,318.1

the number of alerts the subject contains. The alert distribution measures how many FINDBUGS alert types the subject program contains. We are interested in subject programs with varied alert types. The calculation for alert distribution is presented in Equation 4.

$$AD_s = \frac{\Sigma\left(\frac{alerts_{subject,type}}{subjects_{type}}\right)}{alerts_{subject}} \quad (4)$$

Boehm and Turner [2] use polar charts (also called radar charts) to provide a visualization of agile and plan-driven risks in a software development project, and use their visualization to determine which development process fits the project characteristics best. Similarly, we can visualize the characteristics of our possible benchmark subjects' polar charts. Each of the characteristics becomes an axis on the polar chart. Figure 2 presents the polar charts for the six selected FAULTBENCH subjects. In Figure 2, the scale of each axis is normalized. The subjects in FAULTBENCH should have different shapes, which are representative of a variety of subject characteristics. Benchmark selection is quantifiable by taking the area of the polar charts. However, when taking the area of polar charts, the order of the six axes matter otherwise the ordering of subjects by area will be affected. We ordered the axes of the polar charts starting clockwise from the top as shown in Figure 2. We reduced the number of subjects to use in the evaluation of the ARM to six by taking the areas of the polar charts, ordering the subjects by area, and taking every other subject starting with the subject having the largest area. Initially, there were seven subjects, but `mflow` had complicated alert open and closure patterns due to interrelated alert types requiring removal from the current version of FAULTBENCH. Table 1 contains the areas of the polar charts for the potential benchmark subjects. Additionally, the six subjects chosen for the benchmark are shaded in grey.

3.5 FAULTBENCH Subject Initialization

After FAULTBENCH subject selection, the remaining task sample data (alert oracle, source code changes, and experimental control prioritizations) are defined.

3.5.1 Alert Oracle

The first author inspected the source code associated with each static analysis alert and determined if the alerts generated by

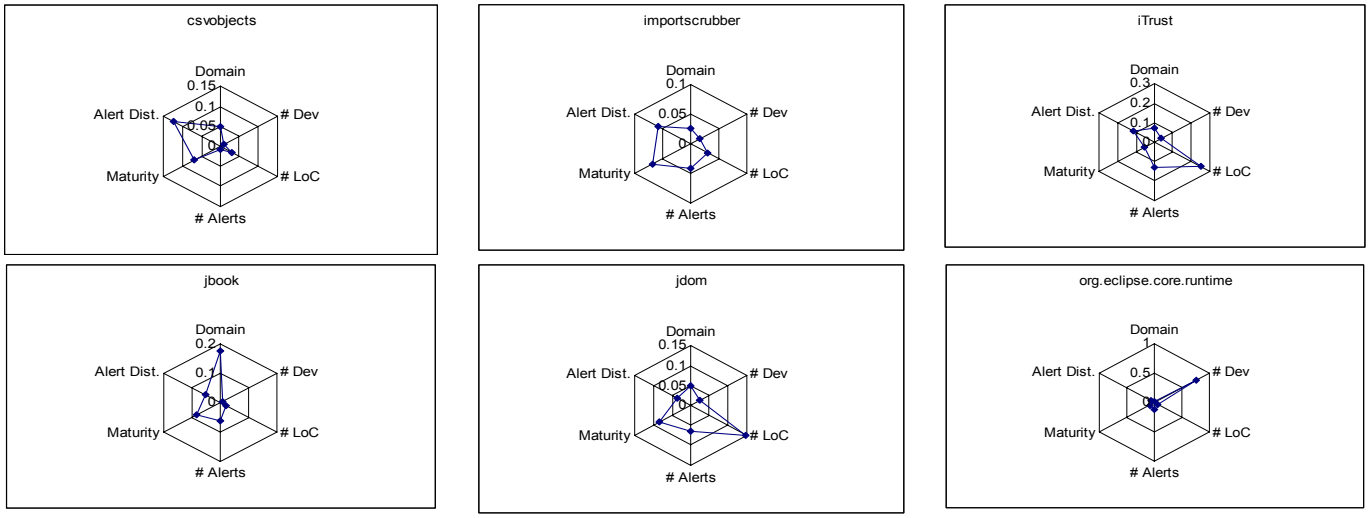

Figure 2: Polar Charts for Potential Benchmark Subjects

FINDBUGS were indications of anomalies in the subject programs. The evaluated alerts provide an oracle for comparing FP mitigation techniques. Table 2 presents the number of TP and FP alerts for each of the benchmark subjects.

Table 2: Benchmark subjects with alert counts

Subject	TP Alerts	FP Alerts	% TP Alerts
csvobjects	3	4	42.9%
importscrubber	11	24	31.4%
iTrust	10	100	9.1%
jbook	26	26	50.0%
jdom	9	46	16.4%
org.eclipse.core.runtime	8	90	8.2%
Average	**11**	**51**	**26.3%**

3.5.2 Source Code Changes

Adaptive FP mitigation techniques modify the prioritization of uninspected alerts from developer feedback about an alert being an indication of an important anomaly or a FP. Modification of an alert's prioritization occurs after each inspection or a set of inspections. An *alert closure* occurs when static analysis tools no longer identify the alert in the source code, usually due to an anomaly fix directly or indirectly associated with the alert description. Alert closures also occur due to configuration changes and file deletions. Alert *suppression* is an explicit action taken by the developer to indicate that an alert is a FP.

The evaluation of adaptive FP mitigation techniques requires fixing TP alerts. However, alerts are not orthogonal; several alert types are interrelated and a change to one may open or close another of a complementary type, which may affect the priority or class of remaining alerts. An alert fix should minimize the number of alert opens and closures. For example, suppose a method opens and closes a file stream within a try block with an empty catch block, as shown in Figure 3. FINDBUGS would report two alerts: 1) OS_OPEN_STREAM_EXCEPTION_PATH at line 3 reporting that the file stream is not closed when there is an exception and 2) DE_MIGHT_IGNORE at line 8 reporting that the exception is ignored. Fixing the alert at line 3, by closing the stream in the exception block, will also close the alert at line 8. When evaluating the prioritization, we only care about closed alerts identified as TP in the baseline. If an alert identified as FP was closed as part of an alert fix, we do not count the alert in our metrics.

Additionally, new alerts may be opened when fixing TP alerts, as shown in the example code in Figure 4. FINDBUGS reports an ES_COMPARING_PARAMETER_STRING_WITH_EQ at line 2, meaning that checking the equality of `a` and a constant string does not use the `.equals()` method. If the alert at line 2 is fixed, a new alert of the same type is opened at line 5 for a similar problem. When inspecting alerts, alerts opened as part of another fix are ignored. Only alerts present in the baseline were inspected and evaluated.

```
1   public void load() {
2     try {
3       BufferedReader in =
4         new BufferedReader(new
5         FileReader(file));
6       //do something with contents
7       in.close();
8     } catch (IOException e) {
9     }
10  }
```

Figure 3: Code Example – Additional Alert Closure

```
1   public void compare(String a) {
2     if (a == "") {
3       //do something
4     }
5     if (a == "null") {
6       //do something else
7     }
8   }
```

Figure 4: Code Example – Alert Opening

3.5.3 Experimental Controls

The OPTIMAL ordering of static analysis alerts has all TP alerts at the top of the alert list; therefore, there are (TP!)(FP!) optimal permutations. For the current version of FAULTBENCH the OPTIMAL ordering is generated by a greedy analysis of the TP alerts. Alerts are initially sorted hierarchically in the context of the subject program (e.g. by project, source folder, class, method, line number, alert type, and description), which provides a repeatable ordering for alerts. To reduce potential bias, prioritization techniques should use the same hierarchical alert ordering to break ties when alerts share the same priority. Alerts are added to the OPTIMAL ordering by the number of TP alerts that are closed when making an alert change. When two alerts close the same number of TP alerts, first the number of FP alerts closed is a tiebreaker, followed by the hierarchical ordering of alerts. At a minimum, the optimal curve will fix one TP alert at each inspection until all TP alerts are fixed.

The TOOL ordering of alerts is created from the tool's alert log information. The RANDOM ordering of alerts is generated via a random number generator[9]. Cases where more than one alert is closed must be considered when creating the OPTIMAL, RANDOM, and TOOL prioritization. The prioritization of an uninspected closed alert is a fraction of the number of alerts closed during an inspection. If there were three alerts (a, b, and c) closed at inspection 3, then the inspected alert (a) would have a rank of 3, the uninspected alert first in the ordered listing (b) would have a rank of 3.33 and other uninspected alert (c) would have a rank of 3.66. For an alert inspected prior to a closure via a tangential change (suppose alert b was inspected at inspection 2), the original inspection (2) is maintained as the rank for that alert (b).

3.6 Benchmark Limitations

The subject programs in FAULTBENCH satisfy the seven desiderata for benchmarks described by Sim et al. [18]. The main limitation of FAULTBENCH is that the first author subjectively chose the classification of the alerts as TP or FP after evaluation of the source code. Another developer may classify alerts differently. Future development of FAULTBENCH by multiple researchers will minimize the subjectivity of the classification. Another limitation is the generation of the OPTIMAL ordering. The current method of generating the OPTIMAL ordering is biased towards the TOOL ordering of alerts. Future work will consider generation of many OPTIMAL orderings such that an average Spearman rank correlation is obtained or the creation of an OPTIMAL ordering of alerts by developers. Finally, the subject programs are all written in Java. Therefore, results obtained on via FAULTBENCH may not be applicable to alert prioritization in other programming languages.

4. BENCHMARK CASE STUDY

We assess the suitability of FAULTBENCH by evaluating three variants of the AWARE-APM [5] FP mitigation technique.

[9] A random sequence generator may be found at http://random.org.

4.1 AWARE-APM

AWARE-APM [5] adaptively prioritizes and classifies static analysis alerts by the likelihood an alert is an indication of an important anomaly. Alerts are prioritized on the continuum, [-1,1] where:

- A priority in [-1,0) implies the alert is likely a FP,
- A priority in (0,1] implies the alert is likely a TP, and
- A priority of 0 means there is not enough information to determine if the alert is likely a TP or FP.

Alerts share *characteristics,* which may demonstrate some causality with the likelihood an alert is a TP. The alert type [5, 11] and alert location [13] are the alert characteristics used in the current version of AWARE-APM, and are presented in Section 4.1.3.

4.1.1 Size Context

The *size context* (SC) represents information about the size of the alerts sharing a characteristic relative to the total number of alerts in a subject. Alerts sharing a characteristic tend to be homogeneous [5, 13], and by increasing the priority of large sets, we can quickly classify many alerts (similar to *information gain* in [13]). The size context is the number of alerts sharing a characteristic divided by the number of alerts for the project. The formula for calculating the size context is presented in Equation 5.

$$SC_c = \frac{\#\,alerts_c}{total\,\#\,alerts} \quad (5)$$

4.1.2 Developer Context

The *developer context* (DC) represents information about what the developer has done to close and suppress alerts while using static analysis during development. We take advantage of homogeneous alert characteristics [5, 13] to utilize the developer's feedback about the alerts to predict the likelihood that other, similar alerts are anomalies. The development context is the difference between closed and suppressed alerts divided by the number of inspected alerts as demonstrated in Equation 6.

$$DC_c = \frac{\#\,\text{closed}_c - \#\,\text{suppressed}_c}{\#\,\text{closed}_c + \#\,\text{suppressed}_c} \quad (6)$$

4.1.3 Alert Characteristics

Alert characteristics are an alert attribute that may have an association with important anomalies. The following subsections describe how we calculate the relationship between alerts sharing the same characteristic like accuracy (ATA) and code locality (CL). The coefficients to the baseline (β_{BC}) and developer (β_{DC}) context have a value of 0.5 implying that the baseline and developer context contribute equally to an alert characteristic calculation.

Alert Type Accuracy (ATA): ATA is the likelihood an alert (a) is an anomaly based on the type of the alert (e.g. null pointer, unclosed stream, etc.) [11, 12]. ATA is the weighted combination of the baseline and developer context of the alert's type. The ATA calculation is described in Equation 7.

$$ATA(a) = \left(\beta_{SC} * SC_{type}\right) + \left(\beta_{DC} * DC_{type}\right) \quad (7)$$

Code Locality (CL): CL is the likelihood an alert (a) is an anomaly based on the location of the alert (e.g. at the source folder, class, or method level). CL is the weighted combination of the baseline and developer context of the alert's location. The contribution of each location is calculated by normalizing the counts of non-singleton source folder, methods, and classes from Table 2b of [13]. The coefficients for the contributions of the source folder, classes, and methods are 0.06, 0.25, and 0.69, respectively and are represented by the coefficients γ_{sf}, γ_c, and γ_m. We are only interested in the non-singleton groups of alerts sharing a characteristic because any action taken on an alert cans be used to predict if the other alerts in the group are likely to be anomalies [13]. Singleton alerts do not provide any predictive data. The calculation for CL is described in Equation 8.

$$CL(a) = \begin{array}{l}\left(\beta_{SC} * \left(\left(\gamma_{sf} * SC_{sf}\right) + \left(\gamma_{sf} * SC_{sf}\right) + \left(\gamma_{sf} * SC_{sf}\right)\right)\right) + \\ \left(\beta_{DC} * \left(\left(\gamma_{sf} * DC_{sf}\right) + \left(\gamma_{sf} * DC_{sf}\right) + \left(\gamma_{sf} * DC_{sf}\right)\right)\right)\end{array} \quad (8)$$

4.1.4 FP Mitigation

The overall alert prioritization calculation is the combination of alert characteristic calculations divided by the number of alert characteristics. Three versions of AWARE-APM FP mitigation techniques are presented in Table 3.

Table 3: Experimental treatments for benchmark evaluation

Treatment	Description or Formula	AWARE Version
ATA	$R(a) = ATA(a)$	1.7.1.1
CL	$R(a) = CL(a)$	1.7.2.0
ATA + CL	$R(a) = \frac{ATA(a) + CL(a)}{2}$	1.7.3.0

4.1.5 ARM Limitations

Similarly to [10, 11, 13], our prioritization technique works best when the groups of alerts sharing a characteristic of interest are fine-grained (e.g. many alert types and locations) and homogeneous. Further research is required to determine how to prioritize static analysis alerts with non-homogeneous groupings.

4.2 Case Study Specifics

Static analysis alerts were prioritized and presented to the developer via the AWARE [5] Eclipse plug-in. AWARE gathers static analysis alerts generated from FINDBUGS and prioritizes the alerts using one of the prioritization functions presented in Table 3. AWARE maintains alert closures and suppressions used to modify the prioritization of the alerts. We used Eclipse version 3.3.1.1 for all of the benchmark subjects except iTrust. For iTrust, we used the Eclipse IDE for Java EE Developers version 3.3.1.1. Each version of AWARE contains one of the three versions of the AWARE-APM FP mitigation techniques. Table 3 also presents the AWARE version for each of the prioritization techniques.

5. CASE STUDY RESULTS

FAULTBENCH provides data to answer the following research questions:

- [Q1]: Can alert prioritization improve the rate of anomaly detection when compared to the tool's output?
- [Q2]: How does the rate of anomaly detection compare between alert prioritization techniques?
- [Q3]: Can alert categorization correctly predict TP and FP alerts?

Question 1 and 2 are answered by using the *area under the curve* metric and the *Spearman rank correlation*, while question 3 is answered using the *precision*, *recall*, and *accuracy* metrics.

5.1 Q1: Improving Anomaly Detection Rate

We plot the cumulative percentage of anomalies detected against the number of inspections and measure the area under the curve to evaluate Question 1. Figure 5 provides an example of these plots for the `jdom` subject program. When TP alerts are fixed, the

Table 4: Area under the fault detection curve for ranking techniques

Subject	Optimal	Random	ATA	CL	ATA + CL	Tool
csvobjects	78.57%	59.52%	50.00%	21.43%	30.39%	54.76%
importscrubber	84.29%	71.82%	66.10%	40.91%	66.62%	36.23%
iTrust	95.5%	48.91%	74.36%	68.09%	67.36%	75.09%
jbook	78.55%	49.83%	46.26%	62.57%	74.19%	39.87%
jdom	91.82%	71.66%	86.16%	63.54%	85.35%	46.89%
org.eclipse.core.runtime	96.81%	68.61%	82.53%	67.09%	82.78%	49.67%
Average	**87.58%**	**61.73%**	**72.57%**	**53.94%**	**67.88%**	**50.42%**

percentage of detected anomaly increases. There are plateaus in the prioritization curve when a FP alert is suppressed at an inspection. A large plateau means there were a number of suppressions. A good prioritization will minimize the large plateaus until most or all of the TP alerts have been identified.

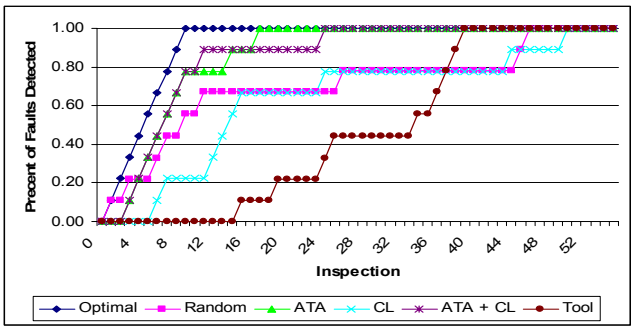

Figure 5: Fault Detection Rate Curves for `jdom`

Table 4 presents the area under the anomaly detection rate curve metrics for each of the prioritization techniques and benchmark subjects. The first question compares alert prioritization techniques to the TOOL ordering of alerts. In the absence of prioritization, developers only have the static analysis tool's output for investigation. If the tool's ordering performs well, then alert prioritization is not needed. However, all prioritization techniques except on `csvobjects` and `iTrust` perform better than the tool ordering. On average, all prioritization techniques have a larger area under the curve (53.94% - 72.57%) than the TOOL ordering (50.42%) of alerts.

Table 5 presents the Spearman rank correlation values between the alert prioritization techniques and OPTIMAL. A positive correlation implies that the specified prioritization is similar to the OPTIMAL prioritization while a negative correlation implies that the specified prioritization is opposite OPTIMAL. The closer the correlation is to 1 or -1, the stronger the match or opposition of the specified prioritization. Cells containing one star (*) have correlations significant at the 0.05 level, while cells containing two stars (**) have correlations significant at the 0.01 level.

Table 5: Spearman Rank Correlation

	ATA	CL	ATA + CL	TOOL
csvobjects	0.321	-0.643	-0.393	0.607
importscrubber	0.512**	-0.026	0.238	0.203
iTrust	0.418**	0.264**	0.261**	0.772**
jbook	0.798**	0.389**	0.599**	-0.002
jdom	0.675**	0.288*	0.457**	0.724**
org.eclipse.core.runtime	0.395**	0.325**	0.246*	0.691**

The TOOL experimental control prioritization has a moderately strong correlation (e.g. correlation value > 0.600) with OPTIMAL for four of the subject programs. The strong correlation is likely due to a similar ordering of the FP alerts, and is not necessarily an indication of the anomaly detection capabilities of the TOOL ordering. For example, the TOOL ordering for `jdom` has a correlation of 0.724; however, the area under the anomaly detection curve for TOOL is at least 20% less than ATA, CL, and ATA+CL as seen in Table 4.

5.2 Q2: Comparing Prioritizations

Table 4 presents the area under the fault detection rate curve metrics for each of the prioritization techniques on FAULTBENCH subjects. The average area under an optimal curve is 90.0%. The ATA prioritization is closer to OPTIMAL than CL prioritization. Additionally, the average ATA area is 30% larger than CL's average area. ATA+CL splits the difference between ATA's and CL's prioritization.

Table 5 presents the Spearman rank correlation values between the alert prioritization technique and OPTIMAL. The correlations between the alert prioritization techniques and OPTIMAL are similar to the patterns observed in the area under the curve measurement in Table 4. However, the ATA correlation with OPTIMAL is typically stronger, indicating that ATA is the better prioritization technique.

5.3 Q3: Categorizing Alerts

Table 6 presents the average precision, recall, and accuracy metrics before each inspection when adaptively categorizing static analysis alerts. We only consider the precision, recall, and accuracy metrics for uninspected alerts because we are trying to predict if the uninspected alerts are TPs or FPs. A priority greater than 0 is a prediction that the alert is a TP while a priority less than 0 is a prediction that an alert is a FP. We then assess the prioritization's classification using the alert oracle and the priority, as shown in Table 7.

If the alert falls in the TP_C or TN_C categories, the prioritization correctly classified the alert as TP or FP. As we learn more about the alerts from the developers, we expect the precision, recall, and accuracy to increase; however, the precision and recall tended to be 0 because after all TP alerts were identified, there was no longer a numerator in the precision and recall equations. The average accuracy is a better measure of how the classification techniques performed. ATA had the best average accuracy, and correctly predicted if an alert is a TP or FP 76% of the time.

Table 7: Alert Classification Assessment

	Alert Oracle	Ranking
True Positive (TP_C)	TP	> 0
True Negative (TN_C)	FP	< 0
False Positive (FP_C)	FP	> 0
False Negative (FN_C)	TP	< 0

Table 6: Average precision, recall, and accuracy metrics of un-inspected alerts at before each inspection

Subject	Average Precision			Average Recall			Average Accuracy		
	ATA	CL	ATA +CL	ATA	CL	ATA +CL	ATA	CL	ATA +CL
csvobjects	0.32	0.50	0.39	.038	.048	0.38	0.58	0.34	0.46
import-scrubber	0.34	0.20	0.18	0.24	0.28	0.45	0.62	0.43	0.56
iTrust	0.05	0.02	0.05	0.16	0.15	0.07	0.97	0.84	0.91
jbook	0.22	0.27	0.23	0.65	0.48	0.61	0.68	0.62	0.66
jdom	0.06	0.09	0.06	0.31	0.07	0.29	0.88	0.86	0.88
org.eclipse.core.runtime	0.05	0.04	0.03	0.17	0.05	0.11	0.92	0.94	0.95
Average	0.17	0.19	0.16	0.42	0.25	0.32	0.76	0.67	0.74

5.4 Benchmark Evaluation

FAULTBENCH contains six programs of varying sizes from several domains. The programs with more than 50 static analysis alerts had more statistically significant results when comparing alert prioritizations with OPTIMAL using the Spearman rank correlation, than the smaller programs. Additionally, if `jbook` or `iTrust` were the only subject used to evaluate alert prioritization techniques the ATA+CL and TOOL prioritizations were the best prioritizations, respectively, when with a larger sample, ATA was the best prioritization technique. The same discrepancy applies when evaluating the classification accuracy of ATA+CL on `org.eclipse.core.runtime`.

The results of the Spearman rank correlation suggest there is bias in the creation of the OPTIMAL order because the TOOL ordering has a moderately strong correlation (> 0.600) with OPTIMAL for four of the subject projects. OPTIMAL defaults to an ordering of alerts by project, source folder, file, method, alert type, line number, and description in the case of a tie. The above ordering is very similar to the TOOL ordering for FINDBUGS due to the use of the Visitor pattern [6]. There are several optimal orderings of alerts, and a semi-randomized ordering may have less bias to the FINDBUGS-TOOL ordering of alerts.

5.5 Case Study Limitations

We consider the three threats to the validity of our case study [17]: construct validity, internal validity, and external validity.

5.5.1 Construct Validity

Construct validity concerns our measurements. The measurements are straight forward and standard for prioritization analysis. In AWARE and the small program used for automating the analysis of the inspection records, possible inconsistencies in our measurements could occur when comparing the static analysis alerts due to line and source code changes during anomaly fix. We consider static analysis alerts to be the same if they share several characteristics including the line number and a hash of the source line. The line number can change through addition or deletion of surrounding code and the source hash can change via refactoring. If both of these characteristics change, we can no longer track the alert. When fixing alerts in the case study, we ensured that only one of the two characteristics was modified for other alerts in the same class. An additional complication is duplicate alerts. An alert is a duplicate when there are two alerts of the same type on the same line of code. The alert display combines the alerts into one listing. Therefore, suppression of the alert listing leads to suppression of both alerts.

5.5.2 Internal Validity

Internal validity concerns the causal relationship between the dependent and independent variables. We are concerned with understanding if the alert characteristics of ATA and CL are indicative of anomalies of importance to the developer. The classification of alerts as TP and FP are from the subjective inspection of the alerts by the first author; therefore, a causal relationship between ATA and CL and the TP and FP classification of the alerts may be exaggerated.

5.5.3 External Validity

External validity concerns how we can generalize our results. Using FAULTBENCH mitigates some of the concerns about generalizing the FP mitigation results due to the varying domains of the subject programs and a larger sample size. Additionally, each of the subject programs is an open source application with real anomalies. However, the programs are relatively small, and there are concerns about scale.

6. CONCLUSIONS AND FUTURE WORK

The literature in the realm of static analysis FP mitigation is moving towards a definition of how to conduct static analysis FP mitigation research [10, 11, 13, 15, 21, 24]. We present FAULTBENCH to supplement the current benchmarks in other languages (e.g. BUGBENCH [15]) and larger Java benchmarks in specific sub-domains (e.g. CHORD subjects [16] for race detection). FAULTBENCH v0.1 is available for use and critique at http://agile.csc.ncsu.edu/faultbench.

We evaluated three alert FP mitigation techniques against the six subjects in FAULTBENCH. Evaluation of the FP mitigation techniques against individual benchmark subjects produced varying results. On `jbook` the ATA+CL prioritization had a larger area under the anomaly detection curve; however, ATA prioritization had a higher rate of anomaly detection on average. In addition, the TOOL ordering performed better than the alert prioritization techniques for `csvobjects`. Individually, the benchmark subjects provide varying results, but together, a larger sample of subject programs provides a better understanding of how FP mitigation techniques work and increase the generalization of experimental conclusions.

We present FAULTBENCH v0.1 to foster collaboration and communication within the static analysis alert prioritization community. We will continue to evolve the benchmark, and feedback from within the community will improve the subjectivity of FAULTBENCH. Additionally, we will continue to investigate static analysis FP mitigation techniques, by analyzing the contributions of the alert characteristic calculations via FAULTBENCH, and modifying the prioritization accordingly.

7. ACKNOWLEDGMENTS

This research is funded by an IBM PhD Fellowship awarded to the first author. We would like to thank the RealSearch reading group, particularly Andy Meneely, and the reviewers for their feedback.

We would like to thank Ben Smith for suggesting the benchmark name.

8. REFERENCES

[1] N. Ayewah, W. Pugh, J. D. Morgenthaler, J. Penix, and Y. Zhou, "Evaluating Static Analysis Defect Warnings On Production Software," *Proceedings of the 7th ACM SIGPLAN-SIGSOFT Workshop on Program Analysis for Software Tools and Engineering*, San Diego, CA, USA, June 13-14, 2007, pp. 1-8.

[2] B. W. Boehm and R. Turner, *Balancing Agility and Discipline: A Guide for the Perplexed*: Addison-Wesley, 2003.

[3] G. Boetticher, T. Menzies, and T. Ostrand, "PROMISE Repository of Empirical Software Engineering Data," http://promisedata.org/ repository, West Virginia University, Department of Computer Science, 2007.

[4] C. Boogerd and L. Moonen, "Prioritizing Software Inspection Results using Static Profiling," *Proceedings of the 6th IEEE Workshop on Source Code Analysis and Manipulation*, Philadelphia, PA, USA, September 27-29, 2006, pp. 149-160.

[5] S. S. Heckman, "Adaptively Ranking Alerts Generated from Automated Static Analysis," in *ACM Crossroads*. vol. 14, no. 1, 2007, pp. 16-20.

[6] D. Hovemeyer and W. Pugh, "Finding Bugs is Easy," *Proceedings of the 19th ACM Conference on Object-Oriented Programming, Systems, Languages, and Applications*, Vancouver, British Columbia, Canada, October 24-28, 2004, pp. 132-136.

[7] M. Hutchins, H. Foster, T. Goradia, and T. Ostrand, "Experiments on the Effectiveness of Dataflow- and Controlflow-Based Test Adequacy Criteria," *Proceedings of the 19th International Conference on Software Engineering*, Sorrento, Italy, May 16-21, 1994, pp. 191-200.

[8] IEEE, "IEEE Standard 610.12-1990, IEEE Standard Glossary of Software Engineering Terminology," 1990.

[9] IEEE, "IEEE 1028-1997 (R2002) IEEE Standard for Software Reviews," 2002.

[10] S. Kim and M. D. Ernst, "Prioritizing Warning Categories by Analyzing Software History," *Proceedings of the International Workshop on Mining Software Repositories*, Minneapolis, MN, USA, May 19-20, 2007, p. 27.

[11] S. Kim and M. D. Ernst, "Which Warnings Should I Fix First?," *Proceedings of the 6th Joint Meeting of the European Software Engineering Conference and the ACM SIGSOFT Symposium on the Foundations of Software Engineering*, Dubrovnik, Croatia, September 3-7, 2007, pp. 45-54.

[12] S. Kim, T. Zimmermann, J. E. James Whitehead, and A. Zeller, "Predicting Faults from Cached History," *Proceedings of the 29th International Conference on Software Engineering*, Minneapolis, MN, USA, May 23-25, 2007, pp. 489-498.

[13] T. Kremenek, K. Ashcraft, J. Yang, and D. Engler, "Correlation Exploitation in Error Ranking," *Proceedings of the 12th ACM SIGSOFT International Symposium on Foundations of Software Engineering*, Newport Beach, CA, USA, 2004, pp. 83-93.

[14] T. Kremenek and D. Engler, "Z-Ranking: Using Statistical Analysis to Counter the Impact of Static Analysis Approximations," *Proceedings of the 10th International Static Analysis Symposium*, San Diego, California, 2002.

[15] S. Lu, Z. Li, F. Oin, L. Tan, P. Zhou, and Y. Zhou, "BugBench: Benchmarks for Evaluating Bug Detection Tools," *Proceedings of the Workshop on the Evaluation of Software Defect Detection Tools*, Chicago, Illinois, 2005.

[16] M. Naik and A. Aiken, "Effective Static Race Detection for Java," *Proceedings of the ACM SIGPLAN 2006 Conference on Programming Language Design and Implementation*, Ottawa, Canada, June 10-16, 2006, pp. 308-319.

[17] G. Rothermel, R. H. Untch, C. Chu, and M. J. Harrold, "Prioritizing Test Cases For Regression Testing," *IEEE Transactions on Software Engineering,* vol. 27, no. 10, pp. 929-948, October, 2001.

[18] S. E. Sim, S. Easterbrook, and R. C. Holt, "Using Benchmarking to Advance Research: A Challenge to Software Engineering," *Proceedings of the 25th International Conference on Software Engineering*, Portland, Oregon, USA, May 3-10, 2003, pp. 74-83.

[19] W. F. Tichy, "Should Computer Scientists Experiment More?," in *Computer*. vol. 31, no. 5, 1998, pp. 32-40.

[20] S. Wagner and M. A. Florian Deissenboeck, Johann Wimmer, Markus Schwalb, "An Evaluation of Two Bug Pattern Tools for Java," *Proceedings of the 1st IEEE International Conference on Software Testing, Verification, and Validation*, Lillehammer, Norway, to appear, 2008.

[21] C. C. Williams and J. K. Hollingsworth, "Automatic Mining of Souce Code Repositories to Improve Bug Finding Techniques," *IEEE Transactions on Software Engineering,* vol. 31, no. 6, pp. 466-480, 2005.

[22] I. H. Witten and E. Frank, *Data Mining: Practical Machine Learning Tools and Techniques*, 2nd ed. Amsterdam: Morgan Kaufmann, 2005.

[23] J. Zheng, L. Williams, N. Nagappan, W. Snipes, J. Hudepohl, and M. Vouk, "On the Value of Static Analysis for Fault Detection in Software," *IEEE Transactions on Software Engineering,* vol. 32, no. 4, pp. 240-253, April, 2006.

[24] T. Zimmermann, R. Premraj, and A. Zeller, "Predicting Defects in Eclipse," *Proceedings of the 3rd International Workshop on Predictor Models in Software Engineering*, Minneapolis, MN, USA, May 20, 2007, p. 9.